Sixtytwo

Reliability infrastructure for on-prem AI training.

When a GPU dies three days into a two-week run, the job should recover instead of restart.

SONNET XU ANTONIO ALONSO-STEPANOV JESSIE DONG

WHAT BREAKS TODAY

On-prem teams still do training reliability by hand.

- A single unhealthy node can kill a multi-day training run.
- Engineers debug across PyTorch, Slurm, Kubernetes, DCGM, and custom scripts.
- Cloud platforms absorb this operational work. On-prem stacks do not.

PROBLEM

Training workloads fail for infrastructure reasons, not model reasons.

What on-prem training teams are dealing with now

01

Long-running jobs

Training runs last for days or weeks. One bad GPU, storage path, or network issue can wipe out expensive compute.

02

Fragmented tooling

The stack spans frameworks, schedulers, metrics, and infra tools that were not built to recover one workflow together.

03

Manual recovery

When something goes wrong, engineers inspect logs, replace nodes, restart jobs, and lose utilization while the cluster waits.

TYPICAL STACK TODAY

PyTorch | Ray | Slurm | Kubernetes | NVIDIA DCGM | Prometheus | OpenTelemetry | Internal scripts

PRODUCT

Monitor the training stack and automate recovery.

Sixtytwo handles the failure path.

01

Detect

Watch GPU telemetry, scheduler state, framework health, and checkpoint status.

02

Checkpoint

Capture model state before the failure cascades across the run.

03

Replace

Quarantine the bad machine and swap in healthy cluster capacity.

04

Resume

Restart the job from a known good point instead of from scratch.

CORE PRODUCT

- Training-aware observability across the full stack.
- Failure detection tied to automated recovery actions.
- Integrations with training frameworks and cluster schedulers.

INITIAL INTEGRATIONS FROM THE APPLICATION

Python | PyTorch | Kubernetes | NVIDIA DCGM | Prometheus | OpenTelemetry

WHY NOW

The workload is moving on-prem before the tooling exists.

Why this category opens now

More owned clusters

At sustained utilization, companies increasingly want to own the hardware instead of paying cloud margins forever.

More custom training

Open weights and proprietary data push teams to train, fine-tune, and distill models in private environments.

Failure cost is rising

The bigger and longer the run, the more expensive each avoidable infrastructure failure becomes.

Tooling gap is still open

Observability tools exist, but they are not built to coordinate checkpointing, node replacement, and resume for AI training.

MARKET

Enterprise software sold to infra and platform teams running training clusters.

Target customers and pricing scale with GPU count

WHO BUYS

- AI companies training large models on owned GPU clusters.
- Enterprises moving sensitive training workloads off the public cloud.
- Infrastructure and platform teams responsible for cluster uptime and utilization.

PRICING MODEL FROM THE YC APPLICATION

Annual enterprise software, priced at roughly **$100 per GPU per month**.

1,000 GPUS
Early large cluster
$1.2M ARR

10,000 GPUS
Scaled deployment
$12M ARR

100,000 GPUS
Category-scale footprint
$120M ARR

COMPETITION

The alternatives are cloud-only, incomplete, or manual.

Competitive landscape

CATEGORY	ON-PREM FOCUS	TRAINING-AWARE	AUTOMATED RECOVERY	UNIFIED STACK VIEW
Cloud training platforms AWS SageMaker / HyperPod	NO	PARTIAL	YES	PARTIAL
DIY stack Slurm + Kubernetes + DCGM + scripts	YES	PARTIAL	NO	NO
Generic observability	PARTIAL	NO	NO	PARTIAL
Sixtytwo	YES	YES	YES	YES

The gap is not raw monitoring. It is reliable recovery for training workloads on owned clusters.

GO TO MARKET

Start with the team that already feels the operational pain.

Wedge, buyer, expansion

WEDGE

Observability + recovery

Start with the narrow pain point: long-running jobs failing and wasting cluster time.

BUYER

Infra / platform teams

Sell to the operators already responsible for uptime, scheduling, and GPU utilization.

EXPANSION

Broader cluster control

Expand into scheduler integrations, fleet health, spare capacity planning, and cluster operations.

CURRENT MOTION FROM THE YC APPLICATION

The team is currently talking to companies running or considering on-prem AI infrastructure, with the goal of shipping an initial prototype in the next few weeks and testing it with a small number of teams running GPU clusters.

CURRENT STATUS

Early stage, pre-launch, with a clear product direction.

Where the company is today

Now

Stage: pre-launch, no revenue yet.

Work underway: customer discovery with teams running or considering on-prem AI clusters.

Near term: initial prototype and pilot testing with a small number of GPU-cluster teams.

Why this team can start here

Training systems: distributed LLM dataloaders and large-scale training work at DatologyAI.

Cluster ops: GPU provisioning, job scheduling, telemetry, and heartbeat recovery work at Mithril.

Production systems: product and ML systems shipped at Meta and other engineering roles.

TEAM

Relevant backgrounds across distributed training, GPU orchestration, and production ML systems.

Founding team

SONNET XU

Training systems

DatologyAI: high-performance dataloader for large-scale distributed LLM training.

Stanford: research background across ML systems and model evaluation.

Why it matters: understands where training throughput and reliability break down.

JESSIE DONG

GPU infrastructure

Mithril: GPU provisioning, job scheduling, telemetry, worker discovery, and heartbeat-based recovery.

Sago: founding engineer building production AI systems.

Why it matters: has worked close to the cluster control plane already.

ANTONIO ALONSO-STEPANOV

Product + systems

Meta: shipped production software including LLM-related internal tooling.

Meta: prior full-stack product engineering internship shipping user-facing systems.

Why it matters: can build the operator product around the infrastructure layer.

The cloud has a reliability layer for training. On-prem needs one too.

Sixtytwo is building the reliability infrastructure for teams training models on owned GPU clusters.

CONTACT

ajalonso@stanford.edu

sonnet@stanford.edu

jessiedong@stanford.edu

ONE-LINE DESCRIPTION

Reliability infrastructure for on-prem AI training.